

Exercises in Praktikum Machine Learning

0. General information

- The goal of this exercise sheet is to implement in MATLAB the extraction of standard features at each pixel from a given image.
- The functions that are computing a feature vector

$$\mathbf{x}(i, j) = [x_1(i, j), \dots, x_{N_{features}}(i, j)]$$

must return a cell array \mathcal{X} of size $1 \times N_{features}$. $\mathcal{X}\{k\}$ is an image such that $[\mathcal{X}\{k\}](i, j)$ is the value of the feature x_k at the location (i, j) .

- The features you are computing must be extracted from a grayscale image of your choice. The value of a given feature over the image can be displayed using the function *imagesc*.
- The size a of patches or kernels must always be an odd number $a = 2\rho + 1$. The centre of the patch/kernel is then clearly defined as the pixel $(\rho + 1, \rho + 1)$.
- Some questions are asked to be solved using a cross-correlation product, generally between an image I of size $n_{rows} \times n_{cols}$ and a kernel K of size $a \times a$ with $a = 2\rho + 1$. As a reminder, the cross-correlation product \star between I and K is the image of size $n_{rows} \times n_{cols}$ defined as

$$[I \star K](i, j) = \sum_{-\rho \leq \alpha, \beta \leq \rho} I(i + \alpha, j + \beta) K(\alpha, \beta)$$

Up to a symmetry applied to the kernel, it is equivalent to a convolution product $*$, defined as

$$[I * K](i, j) = \sum_{-\rho \leq \alpha, \beta \leq \rho} I(i - \alpha, j - \beta) K(\alpha, \beta)$$

We recommend to perform cross-correlation products using the syntax

$$\text{imfilter}(I, K, 'replicate', 'same')$$

which automatically pads the image I on the boundaries.

1. Bank of filters

This section is dedicated to the extraction of features that can be seen as filters over the image. The value of such a feature $x_k(i, j)$ is a linear combination of the intensities within the neighbourhood of (i, j) . More formally, we can write

$$x_k(i, j) = \sum_{-\rho \leq \alpha, \beta \leq \rho} I(i + \alpha, j + \beta) K(\alpha, \beta)$$

where K , called the *cross-correlation kernel*, defines the coefficients of this linear combination.

a) Implement a function *standard_filters.m* taking as parameters:

- the image I
- a cell array \mathcal{F} of strings
- the side a of the square used as kernel (used for ‘mean’ and ‘std’)
- a standard deviation σ (used for ‘gaussian’ and ‘LoG’)

that returns the cell array \mathcal{X} giving the feature responses for the filters specified in the cell array \mathcal{F} . The user must be able to choose one or several strings among the following:

- ‘straight derivatives’ : gives the derivative along rows and columns at each pixel. As an example, the derivation along rows can be defined as:

$$\frac{\partial I}{\partial i}(i, j) = \frac{1}{2}(I(i + 1, j) - I(i - 1, j))$$

- ‘diagonal derivatives’ : gives the derivative along the two diagonals at each pixel.
- ‘mean’ : gives the mean of intensities over the patch of size a .
- ‘std’ : gives the standard deviation of intensities over the patch of size a . Consider applying the mean filter to the squared image.
- ‘gaussian’ : Gaussian kernel of standard deviation σ . You can choose $6\sigma + 1$ as the size of the patch.
- ‘LoG’ : Laplacian-of-Gaussian kernel of standard deviation σ . You can choose $6\sigma + 1$ as the size of the patch.
- ‘all’ : returns all the previous filters for the given a and σ .

Some predefined kernels can be found using *fspecial*. The cell array \mathcal{X} contains at most 8 components (if ‘all’ is selected).

2. Features based on integral images

Integral images are an elegant and fast way to compute the sum (or the mean) of intensities over a rectangle. This section is dedicated to the computation of features based on this principle.

a) From an image I of size $n \times p$, we can define the integral image \tilde{I} of size $(n + 1) \times (p + 1)$ by

$$\begin{aligned} \tilde{I}(i, 1) &= 0 \text{ for } i \in \{1, \dots, n + 1\} \\ \tilde{I}(1, j) &= 0 \text{ for } j \in \{1, \dots, p + 1\} \\ \tilde{I}(i, j) &= \sum_{\substack{i' < i \\ j' < j}} I(i', j') \text{ for } i > 1 \text{ and } j > 1 \end{aligned}$$

Implement a function *integral_image.m* taking as input parameter an image I and returning its integral image \tilde{I} . You can use for this the following identity (valid for $i > 1$ and $j > 1$)

$$\tilde{I}(i, j) = I(i-1, j-1) + \tilde{I}(i-1, j) + \tilde{I}(i, j-1) - \tilde{I}(i-1, j-1)$$

b) Implement a function *mean_patch.m* taking as input arguments

- an integral image \tilde{I}
- the side of a patch a
- the coordinates (α, β) of this patch

that returns the mean of intensities of I over the patch of side a centered on (α, β) . For this, you can notice that, in the configuration of the Figure 1, the sum of intensities of I over the gray rectangle is given by

$$\sum_{\substack{x_1 \leq x \leq x_2 \\ y_1 \leq y \leq y_2}} I(y, x) = \tilde{I}(y_2 + 1, x_2 + 1) - \tilde{I}(y_2 + 1, x_1) - \tilde{I}(y_1, x_2 + 1) + \tilde{I}(y_1, x_1)$$




Figure 1: The integral image \tilde{I} allows a very fast computation of the sum over the gray rectangle of the intensities of I

c) We take a patch of side a centered on the pixel (i, j) of interest and we consider its 8 neighbouring patches as described in Figure 2. By using *mean_patch.m* and a loop over the patch indexes, implement a function *mean_features.m* taking as input arguments

- an integral image \tilde{I}
- the side of the patches a

that returns the feature vector

$$\mathbf{x}(i, j) = [\mu_1(i, j), \mu_2(i, j), \dots, \mu_9(i, j)]$$

as a cell array \mathcal{X} of size 1×9 where $\mu_n(i, j)$ is the average of the intensities over the patch $P_n(i, j)$.

d) With a similar approach, implement a function *lbp.m* (for Local Binary Patterns) taking as input arguments

- an integral image \tilde{I}
- the side of the patches a

that extracts the binary feature vector

$$\mathbf{x}(i, j) = [x_1(i, j), \dots, x_4(i, j), x_6(i, j), \dots, x_9(i, j)]$$

as a cell array \mathcal{X} of size 1×8 , where $x_n(i, j) = 1$ if $\mu_n(i, j) \geq \mu_5(i, j)$ and $x_n(i, j) = 0$ if $\mu_n(i, j) < \mu_5(i, j)$.

e) We propose to compute now the same kind of features but on longer range.

(i) Implement a function *long_range_offset.m* taking as input arguments

- an integral image \tilde{I}
- a side of patches a
- an offset vector $\mathbf{w} = [u, v]$

that returns the feature vector $\mathbf{x}(i, j) = [x_1(i, j), x_2(i, j)]$ as a cell array \mathcal{X} of size 1×2 where:

- $x_1(i, j) = \mu_w(i, j) - \mu_5(i, j)$, where $\mu_w(i, j)$ is the mean over the patch of side a centered on the pixel $(i + u, j + v)$
- $x_2(i, j)$ is the binarised version of $x_1(i, j)$: if $x_1(i, j) \geq 0$ then $x_2(i, j) = 1$, else $x_2(i, j) = 0$

(ii) Implement a similar function *long_range_two_offsets.m* taking as input arguments

- an integral image \tilde{I}
- a side of patches a
- a first offset vector $\mathbf{w}_1 = [u_1, v_1]$
- a second offset vector $\mathbf{w}_2 = [u_2, v_2]$

that returns the difference between the means of the patches centered on $(i + u_1, j + v_1)$ and on $(i + u_2, j + v_2)$, and its binary version. Again, the output is a cell array \mathcal{X} of size 1×2 .




Figure 2: A patch of size $a = 3$ centered on the pixel (i, j) of interest and its 8 neighbouring patches